



A friendly user guide to ELF micropatching with Shiva (alpha v1)

Shiva Manual

1 Introduction to Shiva.....	2
ELF program interpreters.....	2
Shiva is an ELF program interpreter.....	2
Systems and arch's that Shiva supports.....	3
2 Installing Shiva-AArch64.....	3
Download Source code.....	3
Build source code.....	3
Build and install libelfmaster.....	3
Build and install Shiva.....	4
Build and install example patches.....	4
3 Writing patches with Shiva.....	5
The Shiva workflow.....	5
ELF Program compilation and linking example.....	5
ELF Symbol interposition.....	7
Patching an .rodata string (Simple first patch).....	7
Building the patch.....	8
Testing our patch: ro_patch.o.....	8
Using shiva-ld to prelink the patch.....	9
Final outcome of patching test_rodata.....	9

1 Introduction to Shiva

ELF program interpreters

Shiva is an ELF microcode-patching technology that works similarly to the well known (*but little understood*) ELF dynamic linker “/lib/ld-linux.so”. Every ELF program that is dynamically linked has a program header segment type called PT_INTERP that describes the path to the *program interpreter*. The program interpreter is a separate ELF program loaded by the kernel that is responsible for helping to build the process runtime image by loading and linking extra modules containing code and data. Traditionally these modules are called shared libraries, and the program interpreter is “/lib/ld-linux.so”. On aarch64 it is “/lib/ld-linux-aarch64.so.1”

Image 1.0: Requesting program interpreter

```
Elf file type is DYN (Shared object file)
Entry point 0x7e0
There are 8 program headers, starting at offset 64

Program Headers:
  Type           Offset             VirtAddr           PhysAddr
                 FileSiz              MemSiz              Flags  Align
  PHDR           0x0000000000000040 0x0000000000000040 0x0000000000000040
                 0x00000000000001c0 0x00000000000001c0  R      0x8
  INTERP         0x0000000000000200 0x0000000000000200 0x0000000000000200
                 0x000000000000001b 0x000000000000001b  R      0x1
    [Requesting program interpreter: /lib/ld-linux-aarch64.so.1]
  LOAD           0x0000000000000000 0x0000000000000000 0x0000000000000000
                 0x0000000000000a98 0x0000000000000a98  R E    0x10000
  LOAD           0x0000000000000d40 0x00000000000010d40 0x00000000000010d40
                 0x00000000000002d0 0x000000000000042e0  RW     0x10000
  DYNAMIC        0x0000000000000d50 0x00000000000010d50 0x00000000000010d50
                 0x0000000000000200 0x0000000000000200  RW     0x8
  NOTE          0x000000000000021c 0x000000000000021c 0x000000000000021c
                 0x0000000000000044 0x0000000000000044  R      0x4
  GNU_STACK     0x0000000000000000 0x0000000000000000 0x0000000000000000
                 0x0000000000000000 0x0000000000000000  RW     0x10
  GNU_RELRO     0x0000000000000d40 0x00000000000010d40 0x00000000000010d40
                 0x00000000000002c0 0x00000000000002c0  R      0x1
```

Just like any other interpreter, such as Python, or Java, the ELF program interpreter is responsible for performing various types of loading, patching and transformation.

Shiva is an ELF program interpreter

Shiva is a newly designed program interpreter who's sole purpose is to load ELF microcode patches that are written in C and compiled into Shiva modules (ELF relocatable objects). Shiva modules are relocatable objects and by nature already contain the ELF meta-data (*sections, symbols relocations*) necessary to re-build much of the process image. Additionally Shiva has introduced some new concepts such as *Transformations* which are an extension to traditional ELF relocations, allowing for complex patching operations such as function splicing.

The core purpose of Shiva is to create a microcode patching system that seamlessly fits into the existing ELF ABI toolchain of compilers, linkers, and loaders. Shiva aims to load and link patches that are written 100% in C, in a way that is natural for developers who may not be reverse engineers but are experts in C. Shiva is in-itself a linker and a loader that works alongside and in conjunction with the existing Dynamic linker "ld-linux.so". To understand more technical details see "shiva_final_desing.pdf" within the documentation directory.

Systems and arch's that Shiva supports

Currently Shiva's AMP tailored-microcode patching system runs on **Linux AArch64**, and supports **AArch64 Linux ELF PIE binaries**. Future support for ARM32, x86_64/x86_32 and other requested architectures. At the present moment Shiva doesn't support ET_EXEC ELF binaries (non-pie) due to time constraints, but it's on the road map.

2 Installing Shiva-AArch64

Download Source code

```
git clone git@github.com:advanced-microcode-patching/shiva
git clone github.com/elfmaster/libelfmaster
```

Build source code

Build and install libelfmaster

Shiva requires a custom libelfmaster branch that has not been merged into main yet. Specifically make sure to use the *aarch64_support* branch.

```
cd libelfmaster
```

```
git checkout aarch64_support
cd ./src
sudo ./make.sh
```

Build and install Shiva

Commands:

```
cd shiva
make
make shiva-ld
make patches
sudo make install
```

Build results:

Shiva is a static ELF executable and installed to “/lib/shiva”. The Shiva modules are typically stored in “/opt/shiva/modules”. The example patches all live within “/git/shiva/modules/aarch64_patches/”.

Build and install example patches

Please see the “shiva/modules/aarch64_patches/” directory. All of the example patches live here. Let us illustrate an example of building and testing one such patch on a program called ‘test_rodata’.

Commands:

```
cd shiva/modules/aarch64_patches/rodata_interposing
make
sudo make install
```

Test the patch by setting the patch path and running “/lib/shiva” manually.

```
$ SHIVA_MODULE_PATH=./ro_patch.o /lib/shiva ./test_rodata
```

Notice that the software prints the new string “The great arcanum”, per the patch. Yet when the program is executed without Shiva it executes with its old behavior.

```
$ ./test_rodata
```

Install the patch permanently

You may also install the patch permanently with the Shiva prelinking tool. This tool updates the PT_INTERP segment from “/lib/ld-linux.so” to “/lib/shiva”, and updates the dynamic segment with several new tags describing the path to the patch that should be loaded and linked at runtime: “/opt/shiva/modules/ro_patch.o”.

```
$ /bin/shiva-ld -e test_rodata -p ro_patch.o -i /lib/shiva -s /opt/shiva/modules -o test_rodata.patched
```

Shiva now installs the patch in memory everytime it's executed

```
$ ./test_rodata.patched
```

The only discernible difference between test_rodata.patched and test_rodata are the PT_INTERP segment modification and the PT_DYNAMIC addition of two new entries describing the patch path. No actual code or data has changed within the ELF binary. Shiva performs all of the patching at runtime.

3 Writing patches with Shiva

The Shiva workflow


Before delving in to the use of Shiva, let us first explain the workflow of what the ELF compilation and linking process looks like before and after Shiva.

ELF Program compilation and linking example

Standard program compilation begins with a compiler like gcc; the source code files (i.e. main.o and foo.o) get linked into a final output ELF executable my_program.

Static linking of objects


Image 3.0: static linking of two ELF relocatable objects



Additionally these ELF relocatable object files contain calls to functions that live within ELF shared libraries. Once the object files (foo.o and main.o) are linked into an ELF executable (my_program), additional Dynamic linking meta-data is created and stored into what is known as the PT_DYNAMIC segment of the program header table. The PT_DYNAMIC segment contains additional meta-data that is to be passed along to the Program interpreter at runtime. The program interpreter (aka: dynamic linker) is specified as a file path in the PT_INTERP segment: "/lib/ld-linux.so". When the executable is run, the kernel loads and executes the interpreter program first: "/lib/ld-linux.so". The interpreter loads and links the needed shared libraries (i.e. libc.so). The program has now been compiled, and linked twice: once statically, and again dynamically.

Dynamically linked executable

Image 3.1: Dynamically link libc.so in at runtime




A Shiva patch can be written in C and compiled as an ELF relocatable object file. In the same sense that "ld-linux.so" loads and links shared libraries (ELF ET_DYN files), Shiva will load and link patch objects (ELF ET_REL files) to re-write the program in memory just before execution.

Dynamically patched executable

Image 3.2: Chained linking with Shiva

Chained dynamic linking: `"/lib/shiva" & "/lib/ld-linux.so`



ELF Symbol interposition

The ELF format uses symbols to denote code and data. Program functions generally live in the `.text` section, and global variables live within the `.data`, `.rodata`, and `.bss` sections. Shiva makes use of the ELF relocation, section, and symbol data to allow patch developers to *interpose* existing symbols; in other words re-link functions and data that have symbols associated with them. As patch developers we will find this as a natural and beautifully convenient way to patch code and data.

Citation 2.0: From “Oracle linkers and libraries guide”

“Interposition can occur when multiple instances of a symbol, having the same name, exist in different dynamic objects that have been loaded into a process. Under the default search model, symbol references are bound to the first definition that is found in the series of dependencies that have been loaded. This first symbol is said to interpose on the other symbols of the same name.”

Patching an `.rodata` string (Simple first patch)

This patching exercise lives in “`modules/aarch64_patches/rodata_interposing`”. In the following example we are going to look at a simple program that prints a constant string “Arcana Technologies”. Constant data is read-only and is usually stored in the `.rodata` section of a program. The `.rodata` section is generally put into text segment, or a read-only segment that is adjacent to the text segment. A shiva module has it’s own `.rodata` section within it’s mapped text segment at runtime. Our goal with the following patch is to replace a read-only string with another by simply re-defining it in C.

Image 2.0: Source code of program rodata_test.c

```
#include <stdio.h>

const char rodata_string[] = "Arcana Technologies";

int main(void)
{
    printf("rodata_string: %s\n", rodata_string);
    exit(0);
}
```

Image 2.1: Source code of patch object: ro_patch.c

```
elfmaster@esoteric-aarch64:~/amp/shiva/modules/aarch64_patches/rodata_interposing$ cat ro_patch.c
const char rodata_string[] = "The Great Arcanum";
```

Building the patch

The patch source code “modules/aarch64_patches/rodata_interposing/ro_patch.c” is a single line patch, redefining the symbol “rodata_string[]”. To build this patch we can simply run the following gcc command to build the patch:

```
$ gcc -mmodel=large -fno-pic -fno-stack-protector -c ro_patch.c -o ro_patch.o
```

The ELF ET_REL object: ro_patch.o describes a symbol with the same exact name as the symbol **rodata_string** in the target executable. Shiva will see this when linking and it will link in the version of the symbol from our patch. This perfectly demonstrates the concept of symbol interposition in Shiva.

Testing our patch: ro_patch.o

In this example we are patching the AArch64 ELF PIE executable “modules/aarch64_patches/rodata_interposing/test_rodata” Let’s first run the program un-patched.

```
$ ./test_rodata
rodata_string: Arcana Technologies
$
```


The program simply prints the string “rodata_string: Arcana technologies”. To see if our patch worked we can invoke Shiva directly to load and link the executable with the patch at runtime. We must specify the path to our patch object with the SHIVA_MODULE_PATH environment variable.

```
$ SHIVA_MODULE_PATH=./ro_patch.o /lib/shiva ./test_rodata
rodata_string: The Great Arcanum
$
```

When executing ./test_rodata with “/lib/shiva” as the primary loader it installs the specified patch, and the new rodata_string[] value “The Great Arcanum” is printed to stdout. When we run ./test_rodata without “/lib/shiva” it does not install the patch and the original string “Arcana Technologies” is printed.

Using shiva-ld to prelink the patch

The reality is that a user doesn’t want to invoke “/lib/shiva” every single time they want to run a program that needs a patch. This is why “/lib/shiva” is designed to be an ELF interpreter like “ld-linux.so”. We can use the Shiva prelinker “/bin/shiva-ld” to install the appropriate meta-data into the program we are patching:

- Update PT_INTERP with the path to “/lib/shiva” (Replacing “/lib/ld-linux.so”).
- Add several custom entries to PT_DYNAMIC describing the patch basename: “ro_patch.o” and the module search path “/opt/shiva/modules”.

The following command will install the Shiva interpreter path and patch information into the ./test_rodata binary.

```
$ shiva-ld -e test_rodata -p ro_patch.o -i /lib/shiva -s /opt/shiva/modules
$ sudo cp ro_patch.o /opt/shiva/modules
```

Final outcome of patching test_rodata

Now everytime the test_rodata executable is ran it will invoke the Shiva interpreter which will in turn load the patch object ro_patch.o.

```
$ ./test_rodata
rodata_string: The Great Arcanum
$
```

We can observe that the program test_rodata is now printing the patched version of rodata_string when it is ran. This patch may appear permanent, but in effect the patch is actually being installed at runtime by “/lib/shiva” everytime the program runs.

